

CprE 308 Laboratory 7: The FAT-12 Filesystem

Department of Electrical and Computer Engineering
Iowa State University

1 Submission

Submit the following items on Canvas:

- A summary of what you learned or found interesting during this lab, no more than two paragraphs, as a PDF **20 pts**
- Your source code for Exercise 3 (the last part), along with a Makefile **80 pts**

2 Introduction

In this lab, you will gain hands-on experience with the FAT-12 file system. In Part I, you will learn to decode the boot sector by hand, and then write a program to do this automatically. In Part II, you will write a program to perform the `ls` function on a FAT-12 file system “image.”

2.1 FAT-12

FAT-12 is one of a series of file systems based around the concept of a File Allocation Table (FAT)—see the lecture slides on file system implementation. The “12” refers to the number of bits used to store an element in the FAT. The original FAT file system (which had 8-bit entries) was developed in the late 1970s. FAT-12 was developed in 1980 and was used in early versions of DOS, and for much longer on floppy disks. It was replaced by FAT-16, and then FAT-32, which was used extensively in Windows up until Windows NT (or Windows 2000 for consumer desktops).

Wikipedia has extensive articles on FAT. See https://en.wikipedia.org/wiki/File_Allocation_Table for an overview and history of the family of file systems, and https://en.wikipedia.org/wiki/Design_of_the_FAT_file_system for detailed information on the design.

2.2 Terminology

Terminology regarding file systems is notoriously inconsistent. Here we will try to use the following terms.

Partition A large contiguous section of the disk; typically one file system is used per partition.

Sector A small chunk of contiguous bytes on disk; the smallest unit that can be read/written by the disk hardware. Traditionally 512 bytes, may be 4k bytes on modern hardware.

Logical sector A sector numbered according to its position in the disk partition (as opposed to the entire disk).

Block A chunk of contiguous bytes on the file system; the smallest unit that is read/written by the OS. Likely an integer number of sectors.

Cluster Another name for file system-level blocks, used particularly with FAT.

2.3 FAT-12 Filesystem Layout

Table 1: Overview of FAT-12 structure.

	Boot Sector	FAT	2nd FAT	Root Directory	Data
Logical sector number	0	1	X	Y	Z

Table 1 shows an overview of the FAT-12 file system structure. FAT-12 filesystems start with logical sector 0, which is a reserved sector known as the *boot sector*. FAT-12 keeps two copies of the FAT for redundancy. In our case, the first FAT starts at logical sector 1, but after that you have to calculate the starting position of sections using information from the boot sector. After the second FAT comes the root directory, and then the file system's data. The formulas for finding the sector numbers for the 2nd FAT, the root directory, and the data area are as follows:

$$X = 1 + \text{sectors}/\text{FAT}$$

$$Y = X + \text{sectors}/\text{FAT}$$

$$Z = Y + \text{roundup512}(32 \cdot \text{directoryEntries}) /* \text{have to start at the next open block} */$$

Note: `roundup512()` rounds up to the next multiple of 512; this is because data must start at a new block, and in our file system image, the block size is the same as the sector size (as in a floppy disk). We multiply the number of directory entries by 32 because each directory entry takes 32 bytes of space.

2.4 Little-endian numbers

Before we proceed with the lab assignment, let's refresh our understanding of little-endian numbers. Numbers in x86-based systems are stored in the little-endian format. This means that the most significant byte comes last. For example, if the bytes "34 12" were stored at offset 0xF0 in the file, then we would know that 0x34 is stored at 0xF0, and 0x12 is stored at 0xF1. If we interpret this as a "short" value (2 bytes) starting at address 0xF0, then the value appears to be 0x3412, but it represents 0x1234. Similarly, if we were storing an int (4 bytes), the bytes would be reverse order, from least to most significant byte.

Suppose that you had an array of short values in memory. What would they look like? The array is stored forward in memory, so we would reverse the first two bytes to get the first short. We would reverse the next two bytes to get the second short, and so on. What about a character array? Because a character is only one byte, there is no reversal.

3 Part I: Decoding the Boot Sector

3.1 Boot Sector Layout

The **boot sector** stores vital information about the filesystem. The boot sector is laid out in the following way, starting at the beginning of the disk (logical sector 0, byte 0):

All values are stored as **unsigned little-endian** numbers unless otherwise specified.

Table 2: The Layout of FAT-12 Boot Sector

Offset	Length	Contents	Display Format
0x00	3	Binary offset of boot loader	hex
0x03	8	Volume Label (ASCII, null padded)	ASCII
0x0B	2	Bytes per logical sector	decimal
0x0D	1	Logical sectors per cluster (block)	decimal
0x0E	2	Number of reserved sectors	decimal
0x10	1	Number of FATs (generally 2)	decimal
0x11	2	Max number of root directory entries	decimal
0x13	2	Number of logical sectors	decimal
0x15	1	Media descriptor	hex
0x16	2	Logical sectors per FAT	decimal
0x18	2	Physical sectors per track	decimal
0x1A	2	Number of heads	decimal
0x1C	2	Number of hidden sectors	decimal

3.2 Exercise 1

Download files `image` and `bytedump.c` from Canvas. The file “image” is a FAT-12 file system image, and “bytedump.c” is a program to read binary values from the image. Table 2 has shown us how the bits are laid out in the FAT-12 boot sector. You are going to decode “image” and find out the information stored in its boot sector. To help get you started, we’re going to decode a few fields manually using the bytedump program.

Compile “bytedump.c”. Given a filename and an offset (in decimal notation), the program will print out 32 bytes of hex, decimal and ASCII values. For example, `./bytedump image 1536` will produce the following output with the given image:

```
$. /bytedump image 1536
Trying to read 32 bytes in "image" starting from offset 1536.
Actually read 32 bytes:
=====
Address  Hex    Decimal  ASCII
0x0600  0x31    49       1
0x0601  0x36    54       6
0x0602  0x53    83       S
0x0603  0x45    69       E
```

```

0x0604    0x43    67      C
0x0605    0x20    32
...
...
0x061e    0x00     0
0x061f    0x00     0

```

The first column is the address (in hexadecimal), and the second is the data at that address (also in hexadecimal). The third column is the decimal value of the data, while the fourth column is the ASCII value of the data.

Using the bytedump program and the offsets in Table 2, find and decode the values for each of the following fields in the boot sector (remember that it starts at offset 0).

Remember that data is stored in the little endian format.

	Hex	Decimal
Bytes per logical sector		
Logical sectors per cluster (block)		
Max number of root directory entries		
Logical sectors per FAT		

NOTE: No submission is needed for this exercise. You can verify your answers with a TA. You will use these values for debugging the next part.

3.3 Exercise 2

Download “bsdump-template.c” and fill in the missing code marked by comment “TODO”. The completed program should be able to read and then print information from the boot sector of the image. The starting offset and size of each field can be found in Table 2.

An example output of the completed program would be:

```
$ ./bsdump image
```

```

                Name:    mkdosfs
        Bytes/Sector:    512
    Sectors/Cluster:    16
    Reserved Sectors:    1
        Number of FATs:    2
    Root Directory entries: 224
        Logical sectors:    2880
    Medium descriptor:    0x00f0
        Sectors/FAT:        1
        Sectors/Track:    18
        Number of heads:    2
    Number of Hidden Sectors: 0

```

4 Part II: Decoding the Root Directory

In this part of the lab, you will write a program to decode the root directory portion of the filesystem image, and print out the information of all the valid root directory entries.

4.1 Root Directory Layout

The root directory is an array of directory entries (the number of such entries is given in the boot sector). Each directory entry is 32 bytes and is laid out according to the table below:

Offset	Length	Description
0x00	8	Filename
0x08	3	Extension
0x0b	1	Attributes
0x0c	10	Reserved
0x16	2	Last Access Time
0x18	2	Last Access Date
0x1a	2	First Cluster
0x1c	4	Size

Table 3: Directory Entry

Notes on interpreting a directory entry:

4.1.1 Invalid Entries

If the filename starts with 0x00, it's not a valid entry. And, all the entries following this entry are invalid. If it starts with 0xE5, it's also an invalid entry, meaning it has been released/deleted. If it starts with anything else, assume that it's a valid file entry.

4.1.2 File Names

In FAT-12, the file naming convention is 8.3. That is, files have 8-character names and 3-character extensions, such as `FILENAME.EXT`. We will not cover the long filename extensions in this lab.

4.1.3 Attributes

A file's attributes include read/write permissions, among other things. There are 8 bits, but we're only interested in a few of them for this lab, shown in Table 4. If a bit is set, it indicates the property is true. 0 is the least significant (rightmost) bit.

Table 4: Attributes used in this lab.

Bit	Attribute
0	Read-only
1	Hidden
2	System
5	Archive

4.1.4 Date and Time

These fields are each 16 bits long. The bits are laid out as follows:

The time is encoded using 5 bits for the hour, 6 bits for the minutes, and 5 bits for the seconds. One catch - since 5 bits is not enough for 60 seconds, only every other second is counted. In other words, if your counter is 16, that indicates a value of 32 seconds.

The date is encoded in a yyyy/mm/dd format. It uses 7 bits for the years since 1980, so add 1980 to the counter to get the current year. There are 4 bits for the month, and 5 bits for the day.

4.2 Exercise 3

Download `fat12ls-template.c` from Canvas, and fill in the missing code marked with “TODO”. The completed program should do something similar to `ls`: print out the filename, attributes, last access time, last access date, and size of the file for each valid entry in the root directory. No need to sort the files—simply list them in disk order.

For function `decodeBootSector`, you can reuse the code from Exercise 2.

Here is an output example of the completed program:

```
File_Name  Attr    Time      Date       Size
16SEC.TXT  RHS      15:22:50   2002/11/06  331
...
...
BIG.LOG    HS       00:00:58   2009/05/03  62559
(R)ead Only (H)idden (S)ystem (A)rchive
```